

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder

df=pd.read_csv("/content/drive/MyDrive/Breast_Cancer.csv")
```

```
df.head()
```

	id	diagnosis	radius_mean	texture_mean
0	842302	M	17.99	10.38
1	842517	M	20.57	17.77
2	84300903	M	19.69	21.25
3	84348301	M	11.42	20.38
4	84358402	M	20.29	14.34

5 rows × 32 columns

```
df.tail()
```

	id	diagnosis	radius_mean	texture_mean
564	926424	M	21.56	22.39
565	926682	M	20.13	28.25
566	926954	M	16.60	28.08
567	927241	M	20.60	29.33
568	92751	B	7.76	24.54

5 rows × 32 columns

```
df.shape
```

(569, 32)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave_points_mean                  569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
```

Files X



```

12 radius_se          569 non-null    float64
13 texture_se         569 non-null    float64
14 perimeter_se       569 non-null    float64
15 area_se            569 non-null    float64
16 smoothness_se      569 non-null    float64
17 compactness_se     569 non-null    float64
18 concavity_se       569 non-null    float64
19 concave_points_se  569 non-null    float64
20 symmetry_se        569 non-null    float64
21 fractal_dimension_se 569 non-null    float64
22 radius_worst       569 non-null    float64
23 texture_worst      569 non-null    float64
24 perimeter_worst    569 non-null    float64
25 area_worst         569 non-null    float64
26 smoothness_worst   569 non-null    float64
27 compactness_worst  569 non-null    float64
28 concavity_worst    569 non-null    float64
29 concave_points_worst 569 non-null    float64
30 symmetry_worst     569 non-null    float64
31 fractal_dimension_worst 569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

```

```
df.describe()
```

	id	radius_mean	texture_mean	pe
count	5.690000e+02	569.000000	569.000000	
mean	3.037183e+07	14.127292	19.289649	
std	1.250206e+08	3.524049	4.301036	
min	8.670000e+03	6.981000	9.710000	
25%	8.692180e+05	11.700000	16.170000	
50%	9.060240e+05	13.370000	18.840000	
75%	8.813129e+06	15.780000	21.800000	
max	9.113205e+08	28.110000	39.280000	

8 rows × 31 columns

```
df.isnull().values.any()
```

False

```
df.isnull().sum()
```

```

id                0
diagnosis         0
radius_mean       0
texture_mean      0
perimeter_mean    0
area_mean         0
smoothness_mean   0
compactness_mean  0
concavity_mean    0
concave_points_mean 0
symmetry_mean     0
fractal_dimension_mean 0
radius_se         0
texture_se        0
perimeter_se      0
area_se           0
smoothness_se     0
compactness_se    0

```

```
concavity_se      0
concave_points_se 0
symmetry_se       0
fractal_dimension_se 0
radius_worst      0
texture_worst     0
perimeter_worst   0
area_worst        0
smoothness_worst  0
compactness_worst 0
concavity_worst   0
concave_points_worst 0
symmetry_worst    0
fractal_dimension_worst 0
dtype: int64
```

df

	id	diagnosis	radius_mean	texture_mean
0	842302	M	17.99	10.3
1	842517	M	20.57	17.7
2	84300903	M	19.69	21.2
3	84348301	M	11.42	20.3
4	84358402	M	20.29	14.3
...
564	926424	M	21.56	22.3
565	926682	M	20.13	28.2
566	926954	M	16.60	28.0
567	927241	M	20.60	29.3
568	92751	B	7.76	24.5

569 rows × 5 columns

```
df.drop(columns="id",axis=1,inplace=True)
df
```


	diagnosis	radius_mean	texture_mean	perimet
0	M	17.99	10.38	
1	M	20.57	17.77	
2	M	19.69	21.25	
3	M	11.42	20.38	
4	M	20.29	14.34	
...	
564	M	21.56	22.39	
565	M	20.13	28.25	
566	M	16.60	28.08	
567	M	20.60	29.33	
568	B	7.76	24.54	

569 rows × 31 columns

```
X = df.drop('diagnosis', axis = 1)
y = df['diagnosis']
```

```
sns.countplot(y,color='violet')
```

<Axes: xlabel='count', ylabel='diagnosis'>



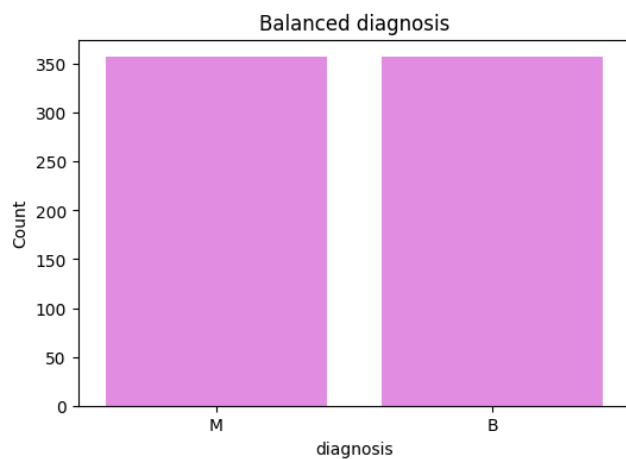
```
import seaborn as sns
import matplotlib.pyplot as plt
from imblearn.over_sampling import RandomOverSampler

# Assuming you have already imported and preprocessed your data

# Upsample the minority class (class 1)
ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(df.drop('diagnosis', axis=1), df['diagnosis'])

# Create a DataFrame with the upsampled data
df_balanced = pd.DataFrame(X_resampled, columns=df.drop('diagnosis', axis=1).columns)
df_balanced['diagnosis'] = y_resampled

# Plot the count of balanced diagnosis after balancing the dataset
plt.figure(figsize=(6, 4))
sns.countplot(x='diagnosis', color='violet', data=df_balanced)
plt.title('Balanced diagnosis')
plt.xlabel('diagnosis')
plt.ylabel('Count')
plt.show()
```



```
from sklearn.preprocessing import LabelEncoder
columns_to_encode = ['diagnosis']
label_encoder = LabelEncoder()
for column in columns_to_encode:
    df[column]=label_encoder.fit_transform(df[column])
cols=df.columns.to_list()
```

df


	diagnosis	radius_mean	texture_mean	perimet
0	1	17.99	10.38	
1	1	20.57	17.77	
2	1	19.69	21.25	
3	1	11.42	20.38	
4	1	20.29	14.34	
...	
564	1	21.56	22.39	
565	1	20.13	28.25	
566	1	16.60	28.08	
567	1	20.60	29.33	
568	0	7.76	24.54	

569 rows × 31 columns

```
X = df.drop('diagnosis', axis = 1)
y = df['diagnosis']
```

```
sns.countplot(df['diagnosis'],color='violet')
```

<Axes: ylabel='count'>



df

	diagnosis	radius_mean	texture_mean	perimet
0	1	17.99	10.38	
1	1	20.57	17.77	
2	1	19.69	21.25	
3	1	11.42	20.38	
4	1	20.29	14.34	
...	
564	1	21.56	22.39	
565	1	20.13	28.25	
566	1	16.60	28.08	
567	1	20.60	29.33	
568	0	7.76	24.54	

569 rows × 5 columns

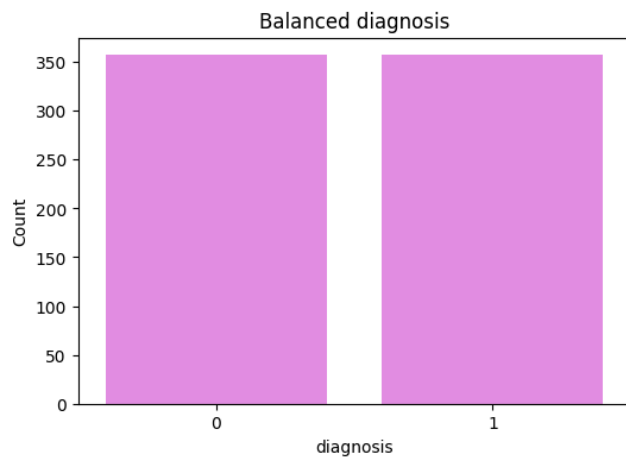
```
import seaborn as sns
import matplotlib.pyplot as plt
from imblearn.over_sampling import RandomOverSampler

# Assuming you have already imported and preprocessed your data

# Upsample the minority class (class 1)
ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(df.drop('diagnosis', axis=1), df['diagnosis'])

# Create a DataFrame with the upsampled data
df_balanced = pd.DataFrame(X_resampled, columns=df.drop('diagnosis', axis=1).columns)
df_balanced['diagnosis'] = y_resampled

# Plot the count of balanced diagnosis after balancing the dataset
plt.figure(figsize=(6, 4))
sns.countplot(x='diagnosis', color='violet', data=df_balanced)
plt.title('Balanced diagnosis')
plt.xlabel('diagnosis')
plt.ylabel('Count')
plt.show()
```



```
# Scaling the data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled

array([[ 1.09706398, -2.07333501,  1.26993369, ...,
        2.29607613,
         2.75062224,  1.93701461],
       [ 1.82982061, -0.35363241,  1.68595471, ...,
        1.0870843 ,
        -0.24388967,  0.28118999],
       [ 1.57988811,  0.45618695,  1.56650313, ...,
        1.95500035,
         1.152255 ,  0.20139121],
       ...,
       [ 0.70228425,  2.0455738 ,  0.67267578, ...,
        0.41406869,
        -1.10454895, -0.31840916],
       [ 1.83834103,  2.33645719,  1.98252415, ...,
        2.28998549,
         1.91908301,  2.21963528],
       [-1.80840125,  1.22179204, -1.81438851, ...,
        -1.74506282,
        -0.04813821, -0.75120669]])
```

```
pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in /usr/local/
Requirement already satisfied: numpy>=1.17.3 in /usr/local/li
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/
Requirement already satisfied: scikit-learn>=1.0.2 in /usr/loc
Requirement already satisfied: joblib>=1.1.1 in /usr/local/li
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/lc
```

```
df
```



	diagnosis	radius_mean	texture_mean	perimet
0	1	17.99	10.38	
1	1	20.57	17.77	
2	1	19.69	21.25	
3	1	11.42	20.38	
4	1	20.29	14.34	
...	
564	1	21.56	22.39	
565	1	20.13	28.25	
566	1	16.60	28.08	
567	1	20.60	29.33	
568	0	7.76	24.54	

569 rows × 31 columns

```
# Oversampling
from imblearn.over_sampling import RandomOverSampler

# Oversample the minority class
oversampler = RandomOverSampler()
X_resampled, y_resampled = oversampler.fit_resample(X_scaled, y)

# Distribution of the target variable 'diagnosis' after oversampling
sns.countplot(y_resampled,color='violet');
```



```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.20,
```

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train_sc = sc.fit_transform(X_train)
X_test_sc = sc.transform(X_test)
```

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
logreg=LogisticRegression()
logreg.fit(X_train,y_train)
y_pred_logreg=logreg.predict(X_test)
from sklearn.metrics import confusion_matrix,accuracy_score,precision_score,recall_score,f1_score
print(confusion_matrix(y_test,y_pred_logreg))
print('Accuracy:',accuracy_score(y_test,y_pred_logreg))
print('Precision:',precision_score(y_test,y_pred_logreg,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_logreg,pos_label=1))
print('F1 score:',f1_score(y_test,y_pred_logreg,pos_label=1))
print(classification_report(y_test, y_pred_logreg))
```

```
[[63  4]
 [ 2 45]]
Accuracy: 0.9473684210526315
Precision: 0.9183673469387755
Recall: 0.9574468085106383
F1 score: 0.9375000000000001
```

	precision	recall	f1-score	support
0	0.97	0.94	0.95	67
1	0.92	0.96	0.94	47
accuracy			0.95	114
macro avg	0.94	0.95	0.95	114
weighted avg	0.95	0.95	0.95	114

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(

KNN

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=7)# default minkowski
knn.fit(X_train,y_train)
y_pred_knn = knn.predict(X_test)
from sklearn.metrics import confusion_matrix,accuracy_score,precision_score,recall_score,f1_score
print(confusion_matrix(y_test,y_pred_knn))
print('Accuracy:',accuracy_score(y_test,y_pred_knn))
print('Precision:',precision_score(y_test,y_pred_knn,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_knn,pos_label=1))
print('f1 score:',f1_score(y_test,y_pred_knn,pos_label=1))
print(classification_report(y_test, y_pred_knn))
```

```
[[64  3]
 [ 3 44]]
Accuracy: 0.9473684210526315
Precision: 0.9361702127659575
Recall: 0.9361702127659575
f1 score: 0.9361702127659575
```

	precision	recall	f1-score	support
0	0.93	0.93	0.93	67
1	0.94	0.94	0.94	47
accuracy			0.94	114
macro avg	0.94	0.94	0.94	114
weighted avg	0.94	0.94	0.94	114

	0	0.96	0.96	0.96	67
	1	0.94	0.94	0.94	47
accuracy				0.95	114
macro avg	0.95	0.95	0.95		114
weighted avg	0.95	0.95	0.95		114

```

from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=7,metric='euclidean')
knn.fit(X_train,y_train)
y_pred_knn=knn.predict(X_test)
from sklearn.metrics import confusion_matrix,accuracy_score,precis
print(confusion_matrix(y_test,y_pred_knn))
print('Accuracy:',accuracy_score(y_test,y_pred_knn))
print('Precision:',precision_score(y_test,y_pred_knn,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_knn,pos_label=1))
print('F1 score:',f1_score(y_test,y_pred_knn,pos_label=1))
print(classification_report(y_test, y_pred_knn))

```

```

[[64  3]
 [ 3 44]]
Accuracy: 0.9473684210526315
Precision: 0.9361702127659575
Recall: 0.9361702127659575
F1 score: 0.9361702127659575

```

	precision	recall	f1-score	support
0	0.96	0.96	0.96	67
1	0.94	0.94	0.94	47
accuracy			0.95	114
macro avg	0.95	0.95	0.95	114
weighted avg	0.95	0.95	0.95	114

```

from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=7,metric='manhattan')
knn.fit(X_train,y_train)
y_pred_knn=knn.predict(X_test)
from sklearn.metrics import confusion_matrix,accuracy_score,precis
print(confusion_matrix(y_test,y_pred_knn))
print('Accuracy:',accuracy_score(y_test,y_pred_knn))
print('Precision:',precision_score(y_test,y_pred_knn,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_knn,pos_label=1))
print('F1 score:',f1_score(y_test,y_pred_knn,pos_label=1))
print(classification_report(y_test, y_pred_knn))

```

```

[[65  2]
 [ 3 44]]
Accuracy: 0.956140350877193
Precision: 0.9565217391304348
Recall: 0.9361702127659575
F1 score: 0.9462365591397849

```

	precision	recall	f1-score	support
0	0.96	0.97	0.96	67
1	0.96	0.94	0.95	47
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

Decision Tree

```

from sklearn.tree import DecisionTreeClassifier
dtree=DecisionTreeClassifier(random_state=1)
dtree.fit(X_train,y_train)
y_pred_dt=dtree.predict(X_test)
from sklearn.metrics import confusion_matrix,accuracy_score,precision_score,recall_score,f1_score
print(confusion_matrix(y_test,y_pred_dt))
print('Accuracy:',accuracy_score(y_test,y_pred_dt))
print('Precision:',precision_score(y_test,y_pred_dt,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_dt,pos_label=1))
print('F1 score:',f1_score(y_test,y_pred_dt,pos_label=1))
print(classification_report(y_test, y_pred_dt))

```

```

[[60  7]
 [ 4 43]]
Accuracy: 0.9035087719298246
Precision: 0.86
Recall: 0.9148936170212766
F1 score: 0.8865979381443299


```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	67
1	0.86	0.91	0.89	47
accuracy			0.90	114
macro avg	0.90	0.91	0.90	114
weighted avg	0.91	0.90	0.90	114

```

from sklearn.tree import plot_tree
fn=list(X_train)
cn=['0','1']
plt.figure(figsize=(4,4),dpi=600)
plot_tree(dtree,feature_names=fn,class_names=cn,filled=True);

```



```

from sklearn.tree import DecisionTreeClassifier
dtree=DecisionTreeClassifier(criterion='entropy')
dtree.fit(X_train,y_train)
y_pred_dt=dtree.predict(X_test)
from sklearn.metrics import confusion_matrix,accuracy_score,precision_score,recall_score,f1_score
print(confusion_matrix(y_test,y_pred_dt))
print('Accuracy:',accuracy_score(y_test,y_pred_dt))
print('Precision:',precision_score(y_test,y_pred_dt,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_dt,pos_label=1))
print('F1 score:',f1_score(y_test,y_pred_dt,pos_label=1))
print(classification_report(y_test, y_pred_dt))

```

```

[[61  6]
 [ 4 43]]

```

```

Accuracy: 0.9122807017543859
Precision: 0.8775510204081632
Recall: 0.9148936170212766
F1 score: 0.8958333333333333

```


```

precision    recall  f1-score   support

```

0	0.94	0.91	0.92	67
1	0.88	0.91	0.90	47
accuracy			0.91	114
macro avg	0.91	0.91	0.91	114
weighted avg	0.91	0.91	0.91	114

```
#Visualizing the tree
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree
fn=list(X_train)
cn=['0','1']
plt.figure(figsize=(4,4),dpi=1000)
plot_tree(dtree,feature_names=fn,class_names=cn,filled=True);
```



```

from sklearn import svm
clf=svm.SVC(kernel='linear',C=0.01)
clf.fit(X_train,y_train)
y_pred_svm=clf.predict(X_test)
from sklearn.metrics import confusion_matrix,accuracy_score,precis
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,y_pred_svm))
print(accuracy_score(y_test,y_pred_svm))
print('Precision:',precision_score(y_test,y_pred_svm,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_svm,pos_label=1))
print('F1 score:',f1_score(y_test,y_pred_svm,pos_label=1))
print(classification_report(y_test,y_pred_svm))

```

```

[[62  5]
 [ 3 44]]
0.9298245614035088
Precision: 0.8979591836734694
Recall: 0.9361702127659575
F1 score: 0.9166666666666666

```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	67
1	0.90	0.94	0.92	47
accuracy			0.93	114
macro avg	0.93	0.93	0.93	114
weighted avg	0.93	0.93	0.93	114

Random Forest-Ensemble Model

```

from sklearn.ensemble import RandomForestClassifier
rf_classifier=RandomForestClassifier(n_estimators=10)
rf_classifier.fit(X_train,y_train)
y_pred_rf=rf_classifier.predict(X_test)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,y_pred_rf))
print(accuracy_score(y_test,y_pred_rf))
print('Precision:',precision_score(y_test,y_pred_rf,pos_label=1))
print('Recall:',recall_score(y_test,y_pred_rf,pos_label=1))
print('F1 score:',f1_score(y_test,y_pred_rf,pos_label=1))
print(classification_report(y_test,y_pred_rf))

```

```

[[66  1]
 [ 3 44]]
0.9649122807017544
Precision: 0.9777777777777777
Recall: 0.9361702127659575
F1 score: 0.9565217391304347

```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	67
1	0.98	0.94	0.96	47
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

Boosting:

Gradient Boosting

AdaBoost

XGBoost

Catboost

```

from sklearn.ensemble import GradientBoostingClassifier
gradient_booster=GradientBoostingClassifier(learning_rate=0.1)
gradient_booster.fit(X_train,y_train)
y_pred_gradboost=gradient_booster.predict(X_test)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,y_pred_gradboost))
print(accuracy_score(y_test,y_pred_gradboost))
print(classification_report(y_test,y_pred_gradboost))

```

```

[[65  2]
 [ 2 45]]
0.9649122807017544

```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	67
1	0.96	0.96	0.96	47
accuracy			0.96	114
macro avg	0.96	0.96	0.96	114
weighted avg	0.96	0.96	0.96	114

```

from sklearn.ensemble import AdaBoostClassifier
abc=AdaBoostClassifier(learning_rate=0.1)
abc.fit(X_train,y_train)
y_pred_abc=abc.predict(X_test)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,y_pred_abc))
print(accuracy_score(y_test,y_pred_abc))
print(classification_report(y_test,y_pred_abc))

```

```

[[65  2]
 [ 3 44]]
0.956140350877193

```

	precision	recall	f1-score	support
0	0.96	0.97	0.96	67
1	0.96	0.94	0.95	47
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

```
pip install xgboost
```

```

Requirement already satisfied: xgboost in /usr/local/lib/python3.10/site-packages (1.7.5)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/site-packages (1.26.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/site-packages (1.11.4)

```



```

from xgboost import XGBClassifier
#print(y_train.unique())
#print(y_test.unique())
#y_train = y_train.replace({'B': 0, 'M': 1})
#y_test = y_test.replace({'B': 0, 'M': 1})
model=XGBClassifier(learning_rate=1)
model.fit(X_train,y_train)
y_pred_xgb=model.predict(X_test)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,y_pred_xgb))
print(accuracy_score(y_test,y_pred_xgb))
print(classification_report(y_test,y_pred_xgb))

```

```

[[67  0]
 [ 2 45]]
0.9824561403508771

```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	67
1	1.00	0.96	0.98	47
accuracy			0.98	114
macro avg	0.99	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

```
pip install catboost
```

```

Collecting catboost
  Downloading catboost-1.2.3-cp310-cp310-manylinux2014_x86_64.
  98.5/98.5 MB 3.3
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: plotly in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: six in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: cyclery>=0.10 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/site-packages (from catboost)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/site-packages (from catboost)
Installing collected packages: catboost
Successfully installed catboost-1.2.3

```

```

from catboost import CatBoostClassifier
model1=CatBoostClassifier()
model1.fit(X_train,y_train)
y_pred_cat=model1.predict(X_test)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,y_pred_cat))
print(accuracy_score(y_test,y_pred_cat))
print(classification_report(y_test,y_pred_cat))

```

```

962: learn: 0.0118500 total: 8.69s remaining:
963: learn: 0.0118732 total: 8.69s remaining:
964: learn: 0.0118458 total: 8.7s remaining:
965: learn: 0.0118232 total: 8.7s remaining:
966: learn: 0.0117976 total: 8.71s remaining:
967: learn: 0.0117850 total: 8.72s remaining:
968: learn: 0.0117577 total: 8.73s remaining:
969: learn: 0.0117296 total: 8.74s remaining:
970: learn: 0.0117062 total: 8.75s remaining:
971: learn: 0.0116965 total: 8.76s remaining:
972: learn: 0.0116789 total: 8.77s remaining:
973: learn: 0.0116427 total: 8.78s remaining:
974: learn: 0.0116143 total: 8.79s remaining:
975: learn: 0.0116113 total: 8.79s remaining:
976: learn: 0.0116029 total: 8.8s remaining:
977: learn: 0.0115812 total: 8.81s remaining:
978: learn: 0.0115531 total: 8.82s remaining:
979: learn: 0.0115317 total: 8.83s remaining:
980: learn: 0.0115307 total: 8.84s remaining:
981: learn: 0.0115189 total: 8.85s remaining:
982: learn: 0.0114960 total: 8.85s remaining:
983: learn: 0.0114938 total: 8.86s remaining:
984: learn: 0.0114851 total: 8.87s remaining:
985: learn: 0.0114655 total: 8.88s remaining:
986: learn: 0.0114497 total: 8.88s remaining:
987: learn: 0.0114241 total: 8.89s remaining:
988: learn: 0.0114160 total: 8.9s remaining:
989: learn: 0.0114035 total: 8.91s remaining:
990: learn: 0.0113811 total: 8.92s remaining:
991: learn: 0.0113660 total: 8.93s remaining:
992: learn: 0.0113450 total: 8.93s remaining:
993: learn: 0.0113286 total: 8.94s remaining:
994: learn: 0.0113075 total: 8.95s remaining:
995: learn: 0.0112741 total: 8.96s remaining:
996: learn: 0.0112649 total: 8.98s remaining:
997: learn: 0.0112396 total: 8.98s remaining:
998: learn: 0.0112154 total: 8.99s remaining:
999: learn: 0.0112034 total: 9s remaining:

```

```
[[66 1]
```

```
[ 2 45]]
```

```
0.9736842105263158
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	67
1	0.98	0.96	0.97	47
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

```

print("y_pred_logreg",y_pred_logreg.shape)
print("y_pred_knn",y_pred_knn.shape)
print("y_pred_svm",y_pred_svm.shape)
print("y_pred_dt",y_pred_dt.shape)

```

```

y_pred_logreg (114,)
y_pred_knn (114,)
y_pred_svm (114,)
y_pred_dt (114,)

```

```
print(X[:5])
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoot
0	17.99	10.38	122.80	1001.0	
1	20.57	17.77	132.90	1326.0	
2	19.69	21.25	130.00	1203.0	
3	11.42	20.38	77.58	386.1	
4	20.29	14.34	135.10	1297.0	

	compactness_mean	concavity_mean	concave_points_mean	symm
0	0.27760	0.3001	0.14710	
1	0.07864	0.0869	0.07017	
2	0.15990	0.1974	0.12790	
3	0.28390	0.2414	0.10520	
4	0.13280	0.1980	0.10430	

	fractal_dimension_mean	...	radius_worst	texture_worst	p
0	0.07871	...	25.38	17.33	
1	0.05667	...	24.99	23.41	
2	0.05999	...	23.57	25.53	
3	0.09744	...	14.91	26.50	
4	0.05883	...	22.54	16.67	

	area_worst	smoothness_worst	compactness_worst	concavity_
0	2019.0	0.1622	0.6656	0.0546
1	1956.0	0.1238	0.1866	0.0460
2	1709.0	0.1444	0.4245	0.0596
3	567.7	0.2098	0.8663	0.0775
4	1575.0	0.1374	0.2050	0.0499

	concave_points_worst	symmetry_worst	fractal_dimension_wor
0	0.2654	0.4601	0.1181
1	0.1860	0.2750	0.0856
2	0.2430	0.3613	0.0870
3	0.2575	0.6638	0.1730
4	0.1625	0.2364	0.0761

[5 rows x 30 columns]

▼ Ensemble

```
from sklearn.ensemble import RandomForestClassifier
X=np.array([y_pred_logreg,y_pred_knn,y_pred_svm,y_pred_dt]).T
meta_learner=RandomForestClassifier()
meta_learner2=meta_learner.fit(X,y_test)
ensemble_prediction=meta_learner.predict(X)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,ensemble_prediction))
print(accuracy_score(y_test,ensemble_prediction))
print(classification_report(y_test,ensemble_prediction))
```

```
[[65  2]
 [ 2 45]]
0.9649122807017544
      precision    recall  f1-score   support

      0       0.97       0.97       0.97         67
      1       0.96       0.96       0.96         47

   accuracy                   0.96         114
  macro avg       0.96       0.96       0.96         114
weighted avg       0.96       0.96       0.96         114
```

```

from sklearn.ensemble import RandomForestClassifier
X=np.array([y_pred_cat,y_pred_xgb,y_pred_abc,y_pred_gradboost,y_pr
meta_learner=RandomForestClassifier()
meta_learner2=meta_learner.fit(X,y_test)
ensemble_prediction=meta_learner.predict(X)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,ensemble_prediction))
print(accuracy_score(y_test,ensemble_prediction))
print(classification_report(y_test,ensemble_prediction))

```

```

[[67  0]
 [ 0 47]]
1.0

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	67
1	1.00	1.00	1.00	47
accuracy			1.00	114
macro avg	1.00	1.00	1.00	114
weighted avg	1.00	1.00	1.00	114

```

from sklearn.ensemble import RandomForestClassifier
X=np.array([y_pred_logreg,y_pred_svm,y_pred_xgb,y_pred_abc,y_pred_
meta_learner=RandomForestClassifier()
meta_learner.fit(X,y_test)
ensemble_prediction=meta_learner.predict(X)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,ensemble_prediction))
print(accuracy_score(y_test,ensemble_prediction))
print(classification_report(y_test,ensemble_prediction))

```

```

[[67  0]
 [ 2 45]]
0.9824561403508771

```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	67
1	1.00	0.96	0.98	47
accuracy			0.98	114
macro avg	0.99	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

```

!pip install lime
import lime
import lime.lime_tabular

```

```

Collecting lime
  Downloading lime-0.2.0.1.tar.gz (275 kB)
    275.7/275.7 kB
  Preparing metadata (setup.py) ... done
Requirement already satisfied: matplotlib in /usr/local/lib/py
Requirement already satisfied: numpy in /usr/local/lib/python3
Requirement already satisfied: scipy in /usr/local/lib/python3
Requirement already satisfied: tqdm in /usr/local/lib/python3.
Requirement already satisfied: scikit-learn>=0.18 in /usr/loc
Requirement already satisfied: scikit-image>=0.12 in /usr/loc
Requirement already satisfied: networkx>=2.2 in /usr/local/li
Requirement already satisfied: pillow!=7.1.0,!>=7.1.1,!>=8.3.0,>
Requirement already satisfied: imageio>=2.4.1 in /usr/local/li
Requirement already satisfied: tifffile>=2019.7.26 in /usr/loc
Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local
Requirement already satisfied: packaging>=20.0 in /usr/local/l
Requirement already satisfied: joblib>=1.1.1 in /usr/local/li

```

```

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/lc
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/
Requirement already satisfied: cyciler>=0.10 in /usr/local/lib/
Requirement already satisfied: fonttools>=4.22.0 in /usr/local
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/
Requirement already satisfied: python-dateutil>=2.7 in /usr/lc
Requirement already satisfied: six>=1.5 in /usr/local/lib/pyt
Building wheels for collected packages: lime
  Building wheel for lime (setup.py) ... done
  Created wheel for lime: filename=lime-0.2.0.1-py3-none-any.v
  Stored in directory: /root/.cache/pip/wheels/fd/a2/af/9ac0a1
Successfully built lime
Installing collected packages: lime
Successfully installed lime-0.2.0.1

```

```

from sklearn import svm
clf=svm.SVC(kernel='linear',C=0.01)
clf.fit(X_train,y_train)
y_pred_svm=clf.predict(X_test)
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(y_test,y_pred_svm))
print(accuracy_score(y_test,y_pred_svm))
print(classification_report(y_test,y_pred_svm))

```

```

[[62  5]
 [ 3 44]]
0.9298245614035088

```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	67
1	0.90	0.94	0.92	47

accuracy			0.93	114
macro avg	0.93	0.93	0.93	114
weighted avg	0.93	0.93	0.93	114

```

[[62  5]
 [ 3 44]]
0.9298245614035088

```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	67
1	0.90	0.94	0.92	47


accuracy			0.93	114
macro avg	0.93	0.93	0.93	114
weighted avg	0.93	0.93	0.93	114

```


from lime.lime_tabular import LimeTabularExplainer
explainer=LimeTabularExplainer(X_train.values,feature_names=X_train
exp=explainer.explain_instance(X_train.values[100],clf.predict,num
exp.as_pyplot_figure()
from matplotlib import pyplot as plt
plt.tight_layout()

```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ba
warnings.warn(
```



```
exp.show_in_notebook(show_table=True)
```




Double-click (or enter) to edit

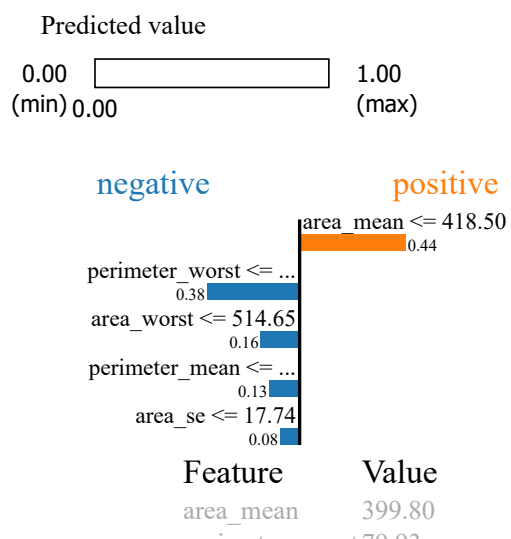
```
from lime.lime_tabular import LimeTabularExplainer
explainer=LimeTabularExplainer(X_train.values,feature_names=X_train
exp=explainer.explain_instance(X_train.values[50],clf.predict,num_
exp.as_pyplot_figure()
from matplotlib import pyplot as plt
plt.tight_layout()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ba
```

```
X does not have valid feature names, but SVC was f
```



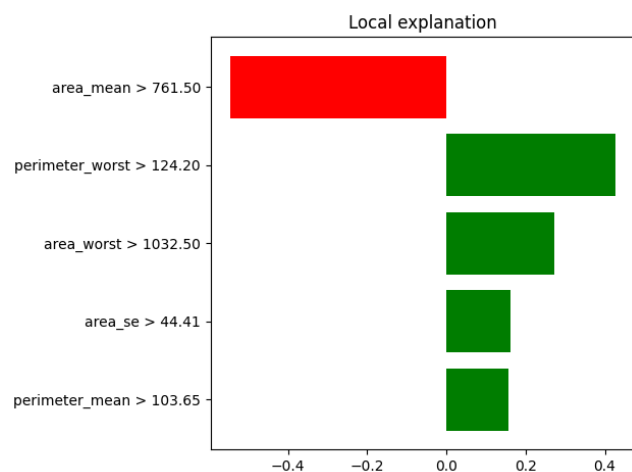
```
exp.show_in_notebook(show_table=True)
```




```
from lime.lime_tabular import LimeTabularExplainer
explainer=LimeTabularExplainer(X_train.values,feature_names=X_train
exp=explainer.explain_instance(X_train.values[100],clf.predict,num
exp.as_pyplot_figure()
from matplotlib import pyplot as plt
plt.tight_layout()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ba
```

```
X does not have valid feature names, but SVC was f
```



```
exp.show_in_notebook(show_table=True)
```




```
from lime.lime_tabular import LimeTabularExplainer
explainer=LimeTabularExplainer(X_train.values,feature_names=X_train
exp=explainer.explain_instance(X_train.values[125],clf.predict,num_
exp.as_pyplot_figure()
from matplotlib import pyplot as plt
plt.tight_layout()
```




```
/usr/local/lib/python3.10/dist-packages/sklearn/ba
```

```
X does not have valid feature names, but SVC was f
```




```
exp.show_in_notebook(show_table=True)
```

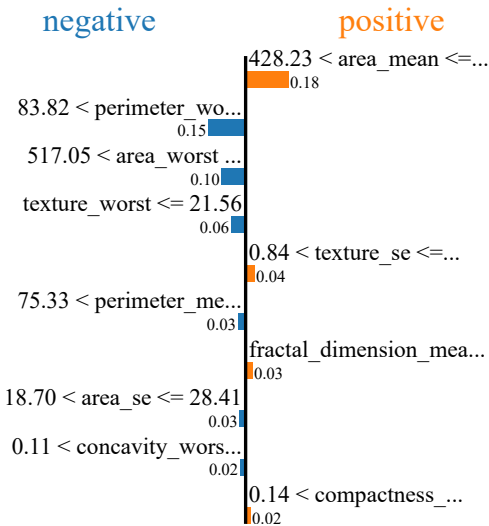



```
from lime.lime_tabular import LimeTabularExplainer
explainer=LimeTabularExplainer(X_test.values,feature_names=X_test.c
exp=explainer.explain_instance(X_test.values[25],clf.predict,num_fe
exp.as_pyplot_figure()
from matplotlib import pyplot as plt
plt.tight_layout()
```

```
warnings.warn(
```



```
exp.show_in_notebook(show_table=True)
```



Feature	Value
area_mean	442.70
perimeter_worst	87.64
area_worst	589.50
texture_worst	20.14
texture_se	1.04
perimeter_mean	76.14
fractal_dimension_mean	0.06

```
pip install shapash
```

```

Requirement already satisfied: Flask<2.3.0 in /usr/local/li
Requirement already satisfied: dash>=2.3.1 in /usr/local/li
Requirement already satisfied: dash-bootstrap-components>=1
Requirement already satisfied: dash-core-components>=2.0.0
Requirement already satisfied: dash-daq>=0.5.0 in /usr/loca
Requirement already satisfied: dash-html-components>=2.0.0
Requirement already satisfied: dash-renderer==1.8.3 in /usr
Requirement already satisfied: dash-table>=5.0.0 in /usr/lo
Requirement already satisfied: nbformat>4.2.0 in /usr/local
Requirement already satisfied: numba>=0.53.1 in /usr/local/
Requirement already satisfied: scikit-learn<1.4,>=1.0.1 in
Requirement already satisfied: category-encoders>=2.6.0 in
Requirement already satisfied: scipy>=0.19.1 in /usr/local/
Requirement already satisfied: statsmodels>=0.9.0 in /usr/l
Requirement already satisfied: patsy>=0.5.1 in /usr/local/l
Requirement already satisfied: Werkzeug<3.1 in /usr/local/l
Requirement already satisfied: importlib-metadata in /usr/l
Requirement already satisfied: typing-extensions>=4.1.1 in
Requirement already satisfied: requests in /usr/local/lib/p
Requirement already satisfied: retrying in /usr/local/lib/p
Requirement already satisfied: nest-asyncio in /usr/local/l
Requirement already satisfied: setuptools in /usr/local/lib
Requirement already satisfied: Jinja2>=3.0 in /usr/local/li
Requirement already satisfied: itsdangerous>=2.0 in /usr/lo
Requirement already satisfied: click>=8.0 in /usr/local/lib
Requirement already satisfied: contourpy>=1.0.1 in /usr/loc
Requirement already satisfied: cycycler>=0.10 in /usr/local/l
Requirement already satisfied: fonttools>=4.22.0 in /usr/lo
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/lo
Requirement already satisfied: packaging>=20.0 in /usr/loca
Requirement already satisfied: pillow>=6.2.0 in /usr/local/
Requirement already satisfied: pyparsing>=2.3.1 in /usr/loc
Requirement already satisfied: python-dateutil>=2.7 in /usr
Requirement already satisfied: fastjsonschema in /usr/local
Requirement already satisfied: jsonschema>=2.6 in /usr/loca
Requirement already satisfied: jupyter-core in /usr/local/l
Requirement already satisfied: traitlets>=5.1 in /usr/local
Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 i
Requirement already satisfied: pytz>=2020.1 in /usr/local/l
Requirement already satisfied: tenacity>=6.2.0 in /usr/loca
Requirement already satisfied: joblib>=1.1.1 in /usr/local/
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/l
Requirement already satisfied: slicer==0.0.7 in /usr/local/
Requirement already satisfied: cloudpickle in /usr/local/li
Requirement already satisfied: MarkupSafe>=2.0 in /usr/loca
Requirement already satisfied: attrs>=22.2.0 in /usr/local/
Requirement already satisfied: jsonschema-specifications>=2
Requirement already satisfied: referencing>=0.28.4 in /usr/
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local
Requirement already satisfied: six in /usr/local/lib/python
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/
Requirement already satisfied: platformdirs>=2.5 in /usr/lo
Requirement already satisfied: charset-normalizer<4,>=2 in
Requirement already satisfied: idna<4,>=2.5 in /usr/local/l
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/l
Requirement already satisfied: certifi>=2017.4.17 in /usr/l

```

```

model=RandomForestClassifier(max_depth=5,random_state=42,n_estimat
model2=model.fit(X_train,y_train)
rf_y_pred=model2.predict(X_test)
rf_y_pred

```

```

array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1,
1, 1, 1, 1,
      0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
0, 0, 1, 0,
      0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1,
0, 0, 1, 0,

```

```



1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0,
1, 0, 0, 0,
1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
1, 0, 1, 1,
0, 1, 1, 0])

```

```

fi=pd.DataFrame({'Feature':X_train.columns,'Importance':model2.feature_importances_})
fi.sort_values(by='Importance',ascending=False,ignore_index=True)

```

	Feature	Importance	
0	concave_points_worst	0.178218	
1	radius_mean	0.108297	
2	concave_points_mean	0.097164	
3	radius_worst	0.094976	
4	perimeter_worst	0.091844	
5	area_worst	0.090185	
6	concavity_worst	0.066257	
7	area_se	0.064503	
8	perimeter_mean	0.050140	
9	compactness_worst	0.034604	
10	symmetry_worst	0.025056	
11	texture_worst	0.013426	
12	texture_se	0.011878	
13	area_mean	0.011625	
14	fractal_dimension_worst	0.011422	
15	perimeter_se	0.007925	
16	concavity_se	0.007664	
17	texture_mean	0.007589	
18	radius_se	0.004198	
19	smoothness_mean	0.003778	
20	symmetry_mean	0.003546	
21	fractal_dimension_mean	0.003253	
22	smoothness_worst	0.003031	
23	fractal_dimension_se	0.002518	
24	symmetry_se	0.002296	
25	smoothness_se	0.001429	
26	compactness_se	0.001406	
27	compactness_mean	0.001284	
28	concavity_mean	0.000490	
29	concave_points_se	0.000000	

```
from shapash.explainer.smart_explainer import SmartExplainer
```

Disk  82.85 GB available 